

DECOMPOSITION & SCHEMA NORMALIZATION

CS 564 - Spring 2025

ACKs: Dan Suciu, Jignesh Patel, AnHai Doan

WHAT IS THIS LECTURE ABOUT?

- Bad schemas lead to redundancy
- To “correct” bad schemas: **decompose** relations
 - lossless-join
 - dependency preserving
- BCNF : a desired **normal forms**

DB DESIGN THEORY

- Helps us identify the “bad” schemas and improve them
 1. express constraints on the data: **functional dependencies (FDs)**
 2. use the FDs to decompose the relations
- The process, called **normalization**, obtains a schema in a “normal form” that guarantees certain properties
 - examples of normal forms: **BCNF, 3NF, ...**

SCHEMA DECOMPOSITION

WHAT IS A DECOMPOSITION?

We **decompose** a relation $\mathbf{R}(A_1, \dots, A_n)$ by creating

- $\mathbf{R}_1(B_1, \dots, B_m)$
- $\mathbf{R}_2(C_1, \dots, C_k)$

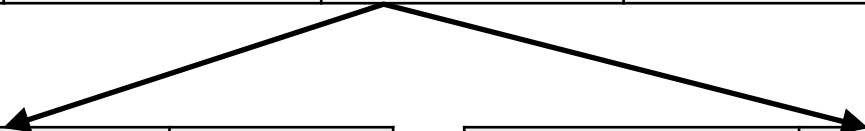
where $\{B_1, \dots, B_m\} \cup \{C_1, \dots, C_k\} = \{A_1, \dots, A_n\}$

- The instance of \mathbf{R}_1 is the projection of \mathbf{R} onto B_1, \dots, B_m
- The instance of \mathbf{R}_2 is the projection of \mathbf{R} onto C_1, \dots, C_k

In general we can decompose a relation into multiple relations.

EXAMPLE: DECOMPOSITION

SSN	name	age	phoneNumber
934729837	Paris	24	608-374-8422
934729837	Paris	24	603-534-8399
123123645	John	30	608-321-1163
384475687	Arun	20	206-473-8221



SSN	name	age
934729837	Paris	24
123123645	John	30
384475687	Arun	20

SSN	phoneNumber
934729837	608-374-8422
934729837	603-534-8399
123123645	608-321-1163
384475687	206-473-8221

DECOMPOSITION DESIDERATA

What should a **good** decomposition achieve?

1. minimize redundancy
2. avoid information loss (**lossless-join**)
3. preserve the FDs (**dependency preserving**)
4. ensure good query performance

EXAMPLE: INFORMATION LOSS

name	age	phoneNumber
Paris	24	608-374-8422
John	24	608-321-1163
Arun	20	206-473-8221

Decompose into:

$R_1(\text{name, age})$

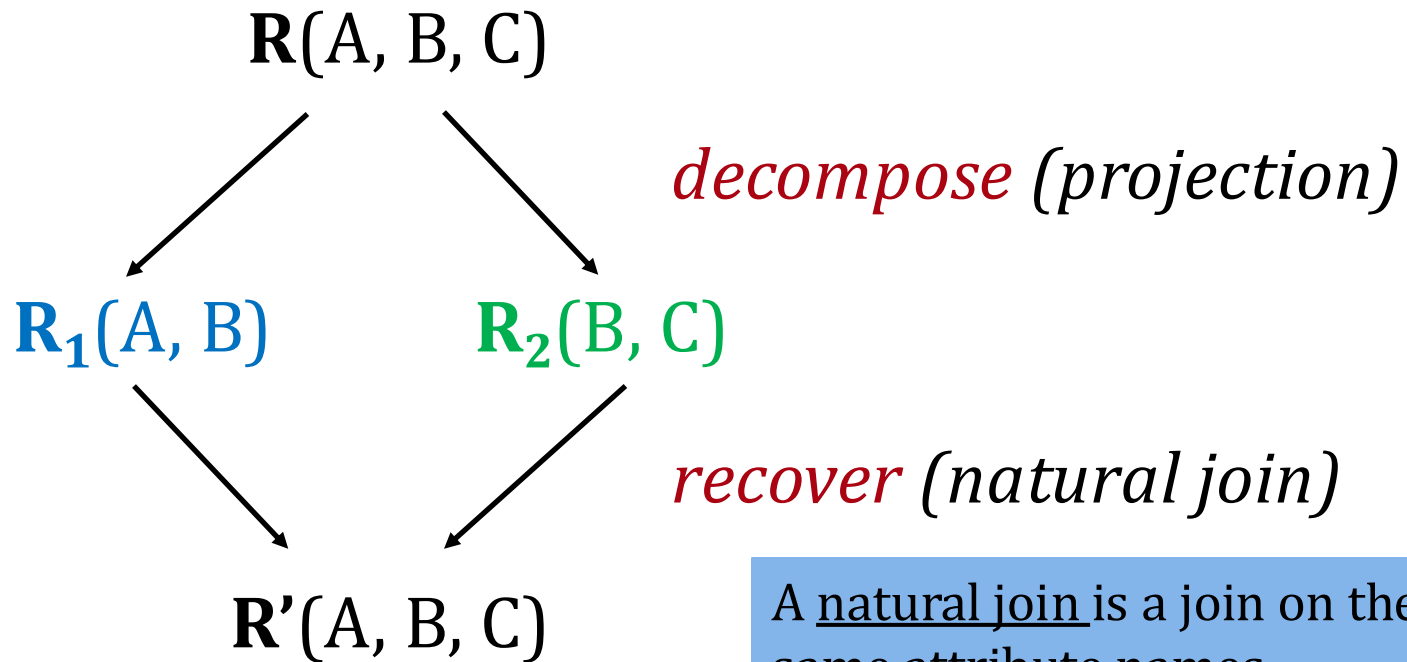
$R_2(\text{age, phoneNumber})$

name	age
Paris	24
John	24
Arun	20

age	phoneNumber
24	608-374-8422
24	608-321-1163
20	206-473-8221

We can't figure out which phoneNumber corresponds to which person!

LOSSLESS-JOIN DECOMPOSITION



A natural join is a join on the same attribute names

A schema decomposition is **lossless-join** if for any initial instance R , $R = R'$

THE CHASE ALGORITHM

The **chase** algorithm is a classic database technique that can be used to check for lossless-join decomposition

Running example

- relation $\mathbf{R}(A, B, C, D, E)$
- FDs: $A \rightarrow B, C$ $D \rightarrow E$

Question: is the following decomposition lossless-join?

$\mathbf{R}_1(A, D)$ $\mathbf{R}_2(A, B, C)$ $\mathbf{R}_3(D, E)$

CHASE: INITIALIZATION

- We create a table with the attributes of the original relation
- We add one row for each relation we split to

	A	B	C	D	E
$R_1(A, D)$	a	b_1	c_1	d	e_1
$R_2(A, B, C)$	a	b	c	d_2	e_2
$R_3(D, E)$	a_3	b_3	c_3	d	e

Attribute not in the relation
gets a subscript

Attribute in the relation - no subscript

CHASE: MAIN ALGORITHM

At every iteration, we check whether an FD is violated, and if so, we “force” it to hold

- If one has a subscript and the other not, we remove the subscript
- If both have a subscript, we make one subscript equal to the other

	A	B	C	D	E
$R_1(A, D)$	a	b_1	c_1	d	$e_1 \rightarrow e$
$R_2(A, B, C)$	a	b	c	d_2	e_2
$R_3(D, E)$	a_3	b_3	c_3	d	e

$A \rightarrow B, C$
 $D \rightarrow E$

The FD $D \rightarrow E$ is violated, so we need to drop the subscript from the first row

CHASE: MAIN ALGORITHM

$A \rightarrow B, C$

$D \rightarrow E$

A	B	C	D	E
a	b_1	c_1	d	e_1
a	b	c	d_2	e_2
a_3	b_3	c_3	d	e

$D \rightarrow E$

A	B	C	D	E
a	b_1	c_1	d	e
a	b	c	d_2	e_2
a_3	b_3	c_3	d	e

$A \rightarrow B, C$

A	B	C	D	E
a	b	c	d	e
a	b	c	d_2	e_2
a_3	b_3	c_3	d	e

At the end of the chase:

- If there is a row without subscripts, we can say that the decomposition is lossless-join
- otherwise, it is not

MORE EXAMPLES

- relation $\mathbf{R}(A, B, C, D)$
- FD $A \rightarrow B, C$

	A	B	C	D
\mathbf{R}_1				
\mathbf{R}_2				

Lossless-join

- decomposition into $\mathbf{R}_1(A, B, C)$ and $\mathbf{R}_2(A, D)$

Not lossless-join

- decomposition into $\mathbf{R}_1(A, B, C)$ and $\mathbf{R}_2(D)$

DEPENDENCY PRESERVING

Given \mathbf{R} and a set of FDs F , we decompose \mathbf{R} into \mathbf{R}_1 and \mathbf{R}_2 . Suppose:

- \mathbf{R}_1 has a set of FDs F_1
- \mathbf{R}_2 has a set of FDs F_2
- F_1 and F_2 are computed from F

A decomposition is dependency preserving if by enforcing F_1 over \mathbf{R}_1 and F_2 over \mathbf{R}_2 , we can enforce F over \mathbf{R}

A NOTE ON FDS OF SPLIT RELATIONS

Given \mathbf{R} and a set of FDs F , we decompose \mathbf{R} into \mathbf{R}_1 and \mathbf{R}_2 . How do we find the FDs F_1 that hold for \mathbf{R}_1 ?

- It is not enough to only keep the FDs from F with attributes in \mathbf{R}_1
- Instead, we need to find the non-trivial FDs in the fd closure of F with attributes in \mathbf{R}_1

Example: $\mathbf{R}(A, B, C)$ with FDs: $A \rightarrow B \quad B \rightarrow C$

- For $\mathbf{R}_1(A, C) \quad F_1 = A \rightarrow C$

GOOD EXAMPLE

Person(SSN, name, age, canDrink)

- $SSN \rightarrow name, age$
- $age \rightarrow canDrink$

decomposes into

- $R_1(SSN, name, age)$
 - $SSN \rightarrow name, age$
- $R_2(age, canDrink)$
 - $age \rightarrow canDrink$

BAD EXAMPLE

$R(A, B, C)$

- $A \rightarrow B$
- $B, C \rightarrow A$

Decomposes into:


- $R_1(A, B)$
 - $A \rightarrow B$
- $R_2(A, C)$
 - no FDs here!!

R_1

A	B
a ₁	b
a ₂	b

R_2

A	C
a ₁	c
a ₂	c



A	B	C
a ₁	b	c
a ₂	b	c

The recovered table
violates $B, C \rightarrow A$

NORMAL FORMS

A **normal form** represents a “good” schema design:

- 1NF (flat tables/atomic values)
- 2NF
- 3NF
- **BCNF**
- 4NF
- ...

more
restrictive



BCNF DECOMPOSITION

BOYCE-CODD NORMAL FORM (BCNF)

A relation **R** is in **BCNF** if whenever $X \rightarrow B$ is a non-trivial FD, then X is a **superkey** in **R**

Equivalent definition: for every attribute set X

- either $X^+ = X$
- or $X^+ = \text{all attributes}$

BCNF EXAMPLE 1

SSN	name	age	phoneNumber
934729837	Paris	24	608-374-8422
934729837	Paris	24	603-534-8399
123123645	John	30	608-321-1163
384475687	Arun	20	206-473-8221

$SSN \rightarrow name, age$

- **key** = $\{SSN, phoneNumber\}$
- $SSN \rightarrow name, age$ is a “bad” FD
- The above relation is **not** in BCNF!

BCNF EXAMPLE 2

SSN	name	age
934729837	Paris	24
123123645	John	30
384475687	Arun	20

$SSN \rightarrow name, age$

- **key** = {*SSN*}
- The above relation is in BCNF!

BCNF EXAMPLE 3

SSN	phoneNumber
934729837	608-374-8422
934729837	603-534-8399
123123645	608-321-1163
384475687	206-473-8221


- **key** = $\{SSN, \textit{phoneNumber}\}$
- The above relation is in BCNF!
- Is it possible that a binary relation is not in BCNF?

BCNF DECOMPOSITION

- Find an FD that violates the BCNF condition

$$A_1, A_2, \dots, A_n \longrightarrow B_1, B_2, \dots, B_m$$

- Decompose **R** to **R₁** and **R₂**:




- Continue until no BCNF violations are left


EXAMPLE

SSN	name	age	phoneNumber
934729837	Paris	24	608-374-8422
934729837	Paris	24	603-534-8399
123123645	John	30	608-321-1163
384475687	Arun	20	206-473-8221

- The FD $SSN \rightarrow name, age$ violates BCNF
- Split into two relations R_1 , R_2 as follows:



EXAMPLE CONT'D



$SSN \rightarrow name, age$

SSN	name	age
934729837	Paris	24
123123645	John	30
384475687	Arun	20

SSN	phoneNumber
934729837	608-374-8422
934729837	603-534-8399
123123645	608-321-1163
384475687	206-473-8221

BCNF DECOMPOSITION PROPERTIES

The BCNF decomposition:

- removes certain types of redundancy
- is **lossless-join**
- is **not always** dependency preserving

BCNF IS LOSSLESS-JOIN

Example:

$R(A, B, C)$ with $A \rightarrow B$ decomposes into:

$R_1(A, B)$ and $R_2(A, C)$

- The BCNF decomposition always satisfies the lossless-join criterion!

BCNF IS NOT DEPENDENCY PRESERVING

$R(A, B, C)$

- $A \rightarrow B$
- $B, C \rightarrow A$

There may not exist any BCNF decomposition that is FD preserving!

The BCNF decomposition is:

- $R_1(A, B)$ with FD $A \rightarrow B$
- $R_2(A, C)$ with no FDs

BCNF EXAMPLE (1)

Books (author, gender, booktitle, genre, price)

- $author \rightarrow gender$
- $booktitle \rightarrow genre, price$

What is the candidate key?

- $(author, booktitle)$ is the only one!

Is it in BCNF?

- **No**, because the left hand side of both (not trivial) FDs is not a superkey!

BCNF EXAMPLE (2)

Books (author, gender, booktitle, genre, price)

- $author \rightarrow gender$
- $booktitle \rightarrow genre, price$

Splitting **Books** using the FD $author \rightarrow gender$:

- **Author** (author, gender)

FD: $author \rightarrow gender$ **in BCNF!**

- **Books2** (author, booktitle, genre, price)

FD: $booktitle \rightarrow genre, price$ **not in BCNF!**

BCNF EXAMPLE (3)

Books (author, gender, booktitle, genre, price)

- $author \rightarrow gender$
- $booktitle \rightarrow genre, price$

Splitting **Books** using the FD $author \rightarrow gender$:

- **Author** (author, gender)
FD: $author \rightarrow gender$ **in BCNF!**
- Splitting **Books2** (author, booktitle, genre, price):
 - **BookInfo** (booktitle, genre, price)
FD: $booktitle \rightarrow genre, price$ **in BCNF!**
 - **BookAuthor** (author, booktitle) **in BCNF!**

IS NORMALIZATION ALWAYS GOOD?

- **Example:** suppose A and B are always used together, but normalization says they should be in different tables
 - decomposition might produce unacceptable performance loss
- **Example:** data warehouses
 - huge historical DBs, rarely updated after creation
 - joins expensive or impractical