

# STORING DATA: DISK AND FILES

---

*CS 564- Spring 2025*

---

*ACKs: Dan Suciu, Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT?

---

How does a DBMS store data?

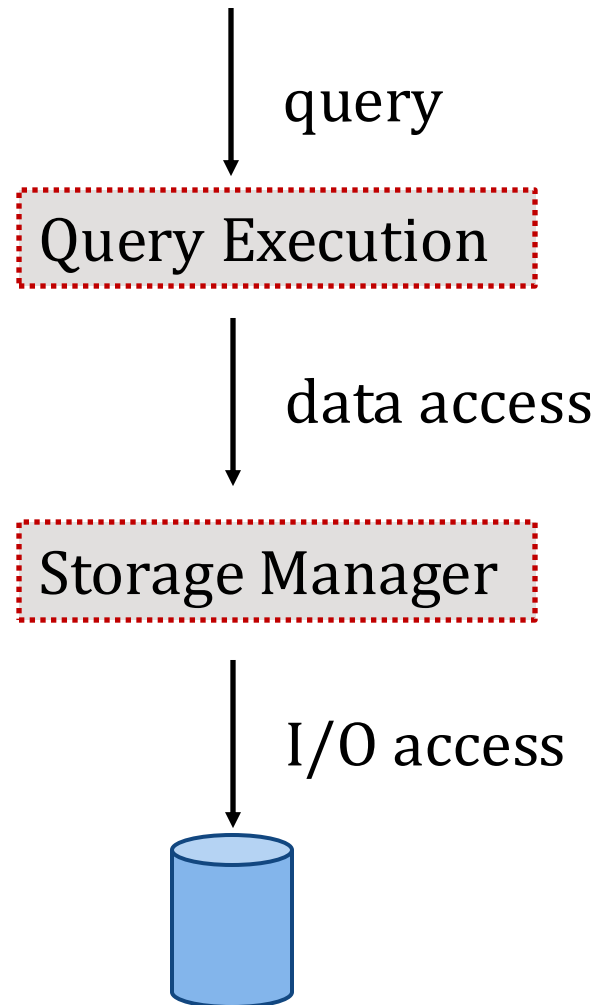
- disk, SSD, main memory

The **buffer manager**

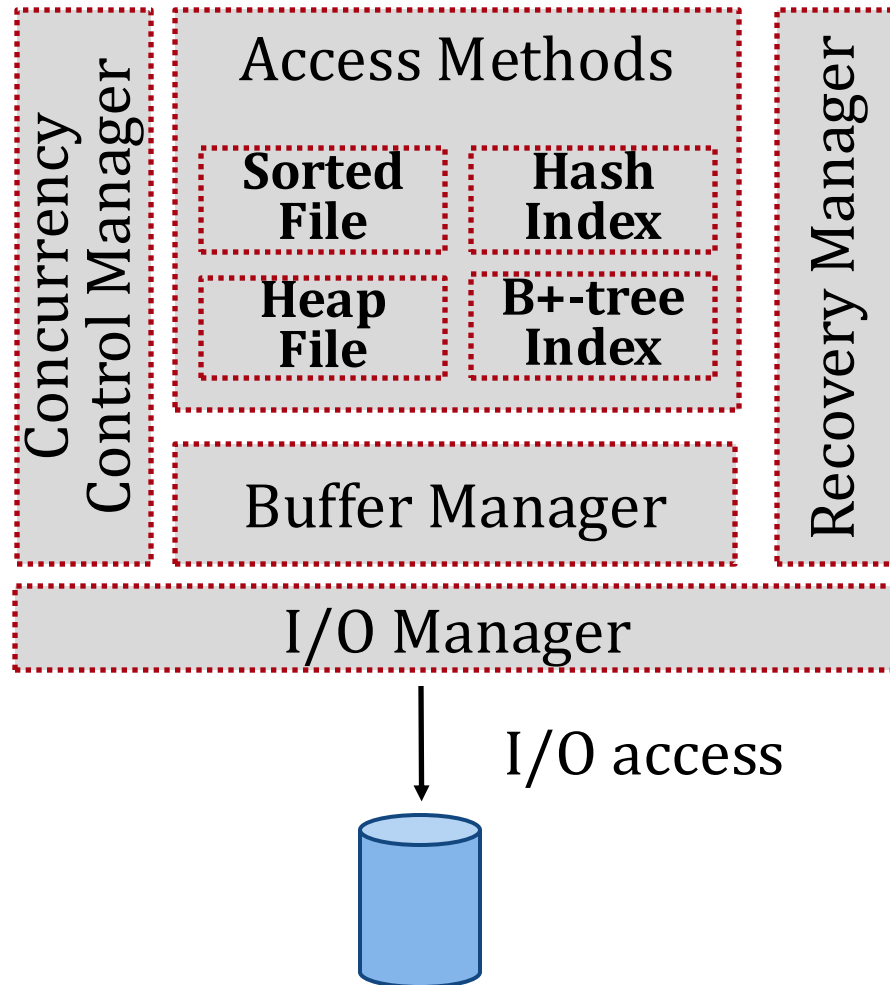
- controls how the data moves between main memory and disk
- uses various replacement policies (LRU, Clock)

# ARCHITECTURE OF A DBMS

---



# ARCHITECTURE OF STORAGE MANAGER




---

# DATA STORAGE

---

- How does a DBMS store and access data?
  - main memory (fast, temporary)
  - disk (slow, permanent)
- How do we move data from disk to main memory?
  - buffer manager
- How do we organize relational data into files?
  - next lecture!

# MEMORY HIERARCHY



---

# WHY NOT MAIN MEMORY?

---

- Relatively high cost
- Main memory is **not persistent!**
- Typical storage hierarchy:
  - **Primary storage:** **main memory** (RAM) for currently used data
  - **Secondary storage:** **disk** for the main database
  - **Tertiary storage:** **tapes** for archiving older versions of the data

---

# DISK

---



---

# DISKS

---

- Secondary storage device of choice
- Data is stored and retrieved in units called disk blocks
- The time to retrieve a disk block varies depending upon location on disk (unlike RAM)


The placement of blocks on disk has major impact on DBMS performance!

# COMPONENTS OF DISKS

- platter: circular hard surface on which data is stored by inducing magnetic changes
- spindle: axis responsible for rotating the platters
- disk head: mechanism to read or write data
- disk arm: moves to position a head on a desired track of the platter


RPM (**R**otations **P**er **M**inute)

- 7200 RPM – 15000 RPM



# COMPONENTS OF DISKS

- data is encoded in concentric circles of sectors called tracks
- block size: multiple of sector size (which is fixed)
- at any time, exactly one head can read/write



---

# ACCESSING THE DISK

---

- unit of read or write: block size
- once in memory, we refer to it as a **page**
- typically: 4k or 8k or 16k

**access time** = **rotational delay** + **seek time** + **transfer time**

# ACCESSING THE DISK (1)

access time = **rotational delay** + **seek time** + **transfer time**

**rotational delay**: time to wait for sector to rotate under the disk head

- typical delay: *0–10 ms*
- maximum delay = 1 full rotation
- average delay  $\sim$  half rotation

RPM	Average delay
5,400	5.56
7,200	4.17
10,000	3.00
15,000	2.00

---

# ACCESSING THE DISK (2)

---

**access time** = **rotational delay** + **seek time** + **transfer time**

**seek time**: time to move the arm to position disk head on the right track

- typical seek time:  $\sim 9\text{ ms}$
- $\sim 4\text{ ms}$  for high-end disks

# ACCESSING THE DISK (3)

**access time** = **rotational delay** + **seek time** + **transfer time**

**data transfer time**: time to move the data to/from the disk surface

- typical rates:  $\sim 100 \text{ MB/s}$
- the access time is dominated by the seek time and rotational delay!

---

# EXAMPLE: SPECS

---

	Seagate HDD
Capacity	<i>3 TB</i>
RPM	<i>7,200</i>
Average Seek Time	<i>9 ms</i>
Max Transfer Rate	<i>210 MB/s</i>
# Platters	<i>3</i>

What are the I/O rates for block size *4 KB* and:

- random workload ( $\sim 0.3 \text{ MB/s}$ )
- sequential workload ( $\sim 210 \text{ MB/s}$ )



---

# EXAMPLE: RANDOM WORKLOAD

---

	Seagate HDD
Capacity	3 TB
RPM	7,200
Average Seek Time	9 ms
Max Transfer Rate	210 MB/s
# Platters	3

For a 4KB block:

- rotational delay = 4.17 ms
- seek time = 9 ms
- transfer time = (4KB) / (210 MB/s)  $\sim$  0.019 ms
- total time per block = 13.1 ms
- I/O rate = (4KB) / (13.1 ms)  $\sim$  0.3 MB/s

---

# ACCESSING THE DISK

---

- Blocks in a file should be arranged sequentially on disk to minimize seek and rotational delay!
- **next** block concept:
  - blocks on same track, followed by
  - blocks on same cylinder, followed by
  - blocks on adjacent cylinder

---

# MANAGING DISK SPACE

---

- The disk space is organized into **files**
- Files are made up of **pages**
- Pages contain **records**
  
- Data is allocated/deallocated in increments of pages
- Logically close pages should be nearby in the disk

# SSD (SOLID STATE DRIVE)

---

- SSDs use flash memory
- **No moving** parts (no rotate/seek motors)
  - eliminates seek time and rotational delay
  - very low power and lightweight
- Data transfer rates: 300-600 MB/s
- SSDs can read data (sequential **or** random) very fast!

# SSDs

---

- Small storage (0.1-0.5x of HDD)
- expensive (20x of HDD)
- **Writes** are much more expensive than **reads** (10x)
- Limited lifetime
  - 1-10K writes per page
  - the average failure rate is 6 years

---

# BUFFER MANAGEMENT

---

---

# BUFFER MANAGER

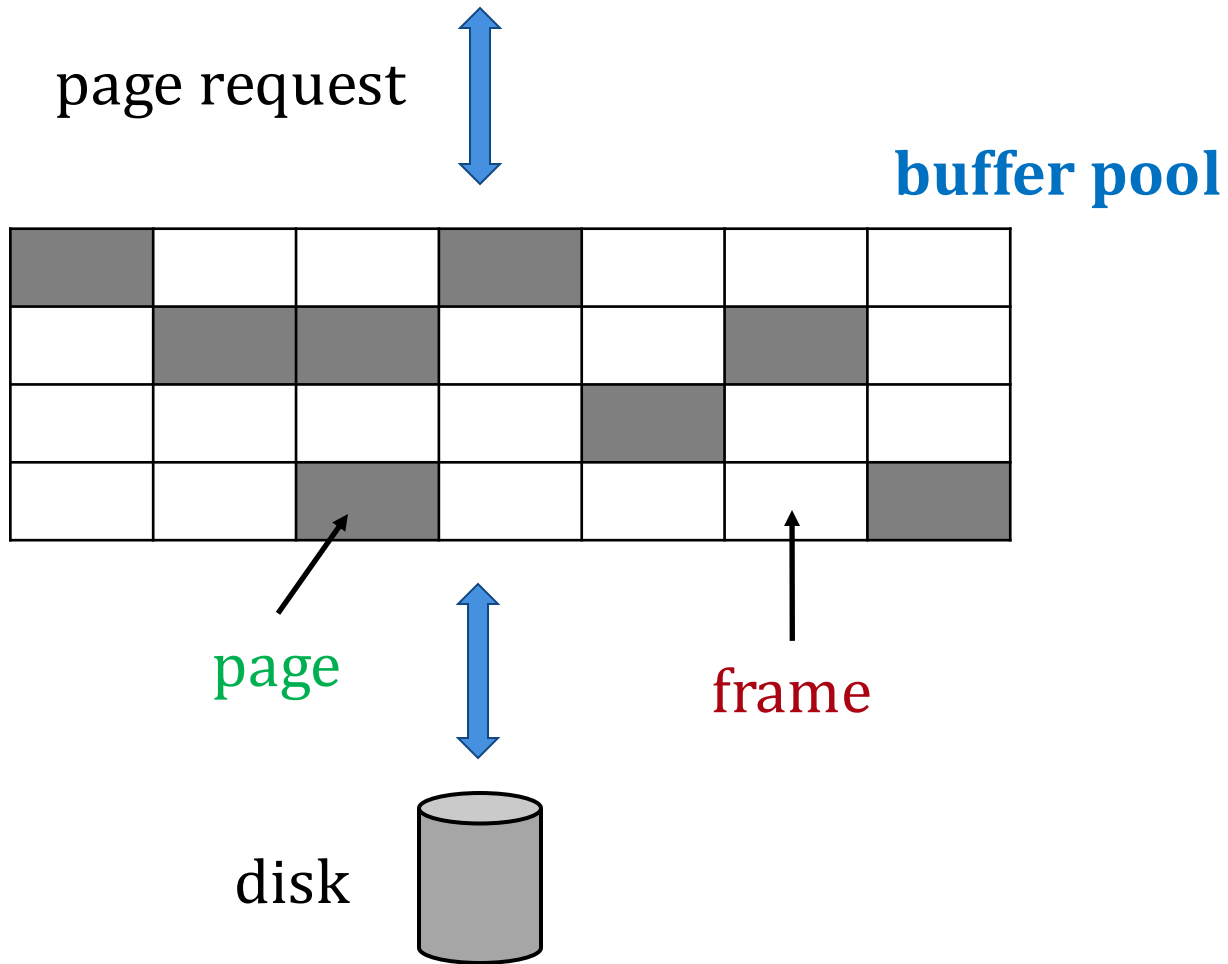
---

- Data must be in RAM for DBMS to operate on it
- All the pages may not fit into main memory

**Buffer manager**: responsible for bringing pages from disk to main memory as needed

- **pages** brought into main memory are in the [buffer pool](#)
- the buffer pool is partitioned into **frames**: slots for holding disk pages

# BUFFER MANAGER





---

# BUFFER MANAGER: REQUESTS

---

- **Read (page)**: read a page from disk and add to the buffer pool (*if not already in buffer*)
- **Flush (page)**: evict page from buffer pool & write to disk
- **Release (page)**: evict page from buffer pool *without* writing to disk

---

# BOOKKEEPING

---

Bookkeeping per frame:

- **pin count**: # current users of the page
  - *pinning* : increment the pin count
  - *unpinning* : decrement the pin count
- **dirty bit**: indicates if the page has been modified
  - **bit = 1** means that the changes to the page must be propagated to the disk

---

# PAGE REQUEST

---

- Page is in the buffer pool:
  - return the address to the frame
  - increment the **pin count**
- Page is not in the buffer pool:
  - choose a frame for replacement (with **pin count = 0**)
  - if frame is **dirty**, write the page to disk
  - read requested page into chosen frame
  - **pin** the page and return the address

---

# BUFFER REPLACEMENT POLICY

---

- How do we choose a frame for replacement?
  - LRU (**L**east **R**ecently **U**sed)
  - Clock
  - MRU (**M**ost **R**ecently **U**sed)
  - FIFO, random, ...
- The replacement policy has big impact on # of I/O's (depends on the access pattern)

# LRU

---

LRU (Least Recently Used)

- uses a **queue** of pointers to frames that have **pin count = 0**
- a page request uses frames only from the *head* of the queue
- when the pin count of a frame goes to 0, it is added to the *end* of the queue

# LRU EXAMPLE

	frame	dirty	pincount
1		0	0
2		0	0
3		0	0

**priority queue:** 1 | 2 | 3

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	0	1
2		0	0
3		0	0

priority queue: 2 | 3

one I/O to read the page

Sequence of requests:

**request A**, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	1
2		0	0
3		0	0

priority queue: 2 | 3

no I/O here!

Sequence of requests:

request A, **modify A**, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E



# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	1
2	B	0	1
3		0	0

priority queue: 3

one I/O to read the page

Sequence of requests:

request A, modify A, **request B**, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	1
2	B	0	2
3		0	0

priority queue: 3

No I/O here  
The pincount increases!

Sequence of requests:

request A, modify A, request B, **request B**, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	2
3		0	0

priority queue: 3 | 1

no I/O yet!

Sequence of requests:

request A, modify A, request B, request B, **release A**,  
request C, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	2
3	C	0	1

priority queue: 1

one I/O to read the page

Sequence of requests:

request A, modify A, request B, request B, release A,  
**request C**, release B, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	1
3	C	0	1

priority queue: 1

the pincount decreases

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, **release B**, request D, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	D	0	1
2	B	0	1
3	C	0	1

priority queue:

two I/Os: one to write A to disk  
and one to read D

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, **request D**, modify D, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	B	0	1
3	C	0	1

priority queue:

no I/O here

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, **modify D**, release B,  
request A, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	B	0	0
3	C	0	1

priority queue: 2

no I/O

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, **release B**,  
request A, request E



# LRU EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	A	0	1
3	C	0	1

priority queue:

one I/O to read A

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
**request A**, request E

# LRU EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	A	0	1
3	C	0	1

priority queue:

The buffer pool is full, the request must wait!

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, **request E**


# CLOCK

---

- Variant of LRU with lower memory overhead
- The  $N$  frames are organized into a cycle
- Each frame has a **referenced bit** that is set to 1 when pin count becomes 0
- A **current** variable points to a frame
- When a frame is considered:
  - If pin count  $> 0$ , increment current
  - If referenced = 1, set to 0 and increment
  - If referenced = 0 and pin count = 0, choose the page

# CLOCK EXAMPLE

	frame	dirty	pincount
1		0	0
2		0	0
3		0	0




Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	0	1
2		0	0
3		0	0




Sequence of requests:

**request A**, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	1
2		0	0
3		0	0



Sequence of requests:

request A, **modify A**, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	1
2	B	0	1
3		0	0

1	0
---	---

3	0
---	---

2	0
---	---

Sequence of requests:

request A, modify A, **request B**, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	1
2	B	0	2
3		0	0

1	0
---	---

3	0
---	---

2	0
---	---

Sequence of requests:

request A, modify A, request B, **request B**, release A,  
request C, release B, request D, modify D, release B,  
request A, request E



# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	2
3		0	0

1	1
---	---

3	0
---	---


2	0
---	---

Sequence of requests:

request A, modify A, request B, request B, **release A**,  
request C, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	2
3	C	0	1




Sequence of requests:

request A, modify A, request B, request B, release A,  
**request C**, release B, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	A	1	0
2	B	0	1
3	C	0	1




Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, **release B**, request D, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	D	0	1
2	B	0	1
3	C	0	1




Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, **request D**, modify D, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	B	0	1
3	C	0	1




Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, **modify D**, release B,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	B	0	0
3	C	0	1



Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, **release B**,  
request A, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	A	0	1
3	C	0	1

1	0
---	---

does a full cycle!

3	0
---	---




2	0
---	---

Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
**request A**, request E

# CLOCK EXAMPLE

	frame	dirty	pincount
1	D	1	1
2	A	0	1
3	C	0	1



Sequence of requests:

request A, modify A, request B, request B, release A,  
request C, release B, request D, modify D, release B,  
request A, **request E**



---

# SEQUENTIAL FLOODING: EXAMPLE

---

- 3 frames in the buffer pool
- request sequence:
  - A, B, C, D, A, B, C, D, A, B, C, D, ...
- With LRU policy, every page access needs an I/O!

	frame	dirty	pincount
1		0	0
2		0	0
3		0	0

---

# SEQUENTIAL FLOODING

---

**Sequential Flooding:** nasty situation caused by LRU policy + repeated sequential scans

- # buffer frames < # pages in file
- each page request causes an I/O !!
- MRU much better in this situation

---

# DBMS vs OS FILE SYSTEM

---

Why not let the OS handle disk management?

- DBMS better at predicting the reference patterns
- Buffer management in DBMS requires ability to:
  - pin a page in buffer pool
  - force a page to disk (for recovery & concurrency)
  - adjust the replacement policy
  - pre-fetch pages based on predictable access patterns
- can better control the overlap of I/O with computation
- can leverage multiple disks more effectively