

BITMAP INDEXES

CS 564- Spring 2025

ACKs: Jignesh Patel, Theo Rekatsinas

WHAT IS THIS LECTURE ABOUT?

- Bitmap Indexes
- Bitslice Indexes

MOTIVATION

Consider the following table:

```
CREATE TABLE Tweets (  
    uniqueMsgID INTEGER,      -- unique message id  
    tstamp      TIMESTAMP,    -- when was the tweet posted  
    uid         INTEGER,      -- unique id of the user  
    msg         VARCHAR (140), -- the actual message  
    zip         INTEGER,      -- zipcode when posted  
    retweet     BOOLEAN       -- retweeted?  
);
```

How can we speed up the following query?

```
SELECT * FROM Tweets  
WHERE zip BETWEEN 53000 AND 54999 ;
```

B+ tree on attribute zip

MOTIVATION

Consider the following table:

```
CREATE TABLE Tweets (  
    uniqueMsgID INTEGER,      -- unique message id  
    tstamp      TIMESTAMP,    -- when was the tweet posted  
    uid         INTEGER,      -- unique id of the user  
    msg         VARCHAR (140), -- the actual message  
    zip         INTEGER,      -- zipcode when posted  
    retweet     BOOLEAN       -- retweeted?  
);
```

How many bytes does a B+ tree use for each record?

- at least key + rid, so: **key-size + rid-size**

Can we do better than that (in terms of storage overhead)?

- **yes!** especially when the attribute domain is small

BITMAP INDEX

THE BITMAP INDEX

- Consider building an index to answer *equality* queries on the **retweet** attribute
- Issues with building a B+ tree:
 - three distinct values: yes, no, NULL
 - many duplicates for each distinct value
 - a weird B+ tree with three long rid lists
- **bitmap index**: build three *bitmap arrays* (stored on disk), one for each value
 - the i^{th} bit in each bitmap corresponds to the i^{th} tuple (we need to map the i^{th} position to a rid!)

BITMAP: EXAMPLE

table (stored in heapfile)

uniqueMsgID	...	zip	retweet
1	...	11324	yes
2	...	53705	yes
3	...	53706	no
4	...	53705	NULL
5	...	90210	no
...
1,000,000,000	...	53705	yes

bitmap index (on retweet)

yes	no	null
1	0	0
1	0	0
0	1	0
0	0	1
0	1	0
...
1	0	0

SELECT * FROM Tweets WHERE retweet = "no" ;

- scan the "no" bitmap file
- for each bit set to 1, compute the tuple rid
- fetch the tuple

A CRITICAL ISSUE

- We need an efficient way to compute a bit position:
 - layout the bitmap in page-id order
- We need an efficient way to map a bit position to a rid:
 - fix the # records per page in the heapfile
 - lay the pages out so that page-ids are sequential and increasing
 - then construct **rid (page-id, slot#)**
 - **page-id** = bit-position / #records-per-page
 - **slot#** = bit-position % #records-per-page

With variable length records, we have to set the limit based on the size of the largest record, which may result in under-filled pages!

BITMAP: OTHER QUERIES

table (stored in heapfile)

uniqueMsgID	...	zip	retweet
1	...	11324	yes
2	...	53705	yes
3	...	53706	no
4	...	53705	NULL
5	...	90210	no
...
1,000,000,000	...	53705	yes

bitmap index (on retweet)

yes	no	null
1	0	0
1	0	0
0	1	0
0	0	1
0	1	0
...
1	0	0

```
SELECT COUNT(*) FROM Tweets WHERE retweet = "no" ;
```

```
SELECT * FROM Tweets WHERE retweet IS NOT NULL ;
```

STORING A BITMAP INDEX

- One bitmap for each value, and one for NULL
- to store each bitmap, use one file for each
- Bitmaps can be compressed!

index size = #tuples * (domain size + 1) *bits*

When is a bitmap more space efficient than a B+ tree?

#distinct values < data entry size in the B+ tree

BITSLICE INDEX

MOTIVATION

Reconsider the following table:

```
CREATE TABLE Tweets (  
    uniqueMsgID INTEGER,      -- unique message id  
    tstamp      TIMESTAMP,    -- when was the tweet posted  
    uid         INTEGER,      -- unique id of the user  
    msg         VARCHAR (140), -- the actual message  
    zip         INTEGER,      -- zipcode when posted  
    retweet     BOOLEAN       -- retweeted?  
);
```

with the following query:

```
SELECT * FROM Tweets WHERE zip = 53706 ;
```

Building a bitmap index on zip is not a good idea!

BITSLICE INDEX

table (stored in heapfile)

uniqueMsgID	...	zip	retweet
1	...	11324	yes
2	...	53705	yes
3	...	53706	no
4	...	53705	NULL
5	...	90210	no
...
1,000,000,000	...	53705	yes

convert to binary

bitslice index

00010110000111100	
01101000111001001	
01101000111001010	
01101000111001001	
10110000001100010	
...	
01101000111001001	

1 slice per bit
+ (possibly) one more slice for NULL

slice 16
higher bit

slice 0
lower bit

BITSLICE INDEX: QUERIES

...	zip
...	11324
...	53705
...	53706
...	53705
...	90210
...	...
...	53705

00010110000111100
01101000111001001
01101000111001010
01101000111001001
10110000001100010
...
01101000111001001
00010111011100000

SELECT * FROM Tweets
WHERE zip <= 12000 ;

1
0
0
0
0
...
0

= **12000** in binary

slice 16

slice 0

walk through each slice constructing a **result bitmap**

- If we look for 0 and have 1, put 0 in the result
- If we look for 1 and have 0, put 1 in the result
- Else we need to consider the next bitslice

OTHER QUERIES

- We can also do **aggregates** with bitslice indices:
 - e.g. `SUM(attr)`: add bitslice by bitslice
 - count the number of 1s in **slice 16** and multiply the count by 2^{16}
 - count the number of 1s in **slice 15** and multiply the count by 2^{15}
 - ...
- We can store each slice using methods like what we have for a bitmap (we can compress again!)

BITMAP VS BITSlice INDEX

- Bitmaps are better for low cardinality domains
- Bitslices are better for high cardinality domains
- It is generally easier to “do the math” with bitmap indices